

Database Query Optimization

1) Existing Query:

```
EXPLAIN ANALYZE
SELECT
    t.employee_id,
    (SELECT name FROM employee e WHERE e.id = t.employee_id),
    (SELECT name FROM project p WHERE p.id = t.project_id),
    SUM(EXTRACT(EPOCH FROM (t.time_to - t.time_from)) / 3600)
FROM time_record t
WHERE t.time_from >= NOW() - INTERVAL '1 month'
GROUP BY
    t.employee_id,
    (SELECT name FROM employee e WHERE e.id = t.employee_id),
    (SELECT name FROM project p WHERE p.id = t.project_id)
ORDER BY
    (SELECT name FROM employee e WHERE e.id = t.employee_id),
    (SELECT name FROM project p WHERE p.id = t.project_id);
```

EXPLAIN ANALYZE result:

```
"GroupAggregate (cost=7450.14..10514.78 rows=187 width=592) (actual
time=0.400..0.403 rows=2 loops=1)"
"  Group Key: ((SubPlan 1)), ((SubPlan 2)), t.employee_id"
"  -> Sort (cost=7450.14..7451.27 rows=453 width=576) (actual time=0.372..0.373
rows=3 loops=1)"
"    Sort Key: ((SubPlan 1)), ((SubPlan 2)), t.employee_id"
"    Sort Method: quicksort Memory: 25kB"
"    -> Seq Scan on time_record t (cost=0.00..7430.16 rows=453 width=576) (actual
time=0.305..0.313 rows=3 loops=1)"
"      Filter: (time_from >= (now() - '1 mon'::interval))"
"      SubPlan 1"
"        -> Index Scan using employee_pkey on employee e (cost=0.15..8.17
rows=1 width=138) (actual time=0.028..0.028 rows=1 loops=3)"
"          Index Cond: (id = t.employee_id)"
"      SubPlan 2"
"        -> Index Scan using project_pkey on project p (cost=0.14..8.16 rows=1
width=418) (actual time=0.021..0.021 rows=1 loops=3)"
"          Index Cond: (id = t.project_id)"
"Planning Time: 0.755 ms"
"Execution Time: 0.629 ms"
```

Identified bottleneck:

- No Join were used in the existing query
- Found there were 2 SubPlans and loops 3 times as per 3 records row to each employee and project in the existing query. If the dataset is larger, then the subquery will loop many times

2) Optimized Query:

```
EXPLAIN ANALYZE
SELECT t.employee_id,
       e.name,
       p.name,
       SUM(EXTRACT(EPOCH FROM (t.time_to - t.time_from)) / 3600)
total_hours
FROM time_record t
INNER JOIN employee e ON e.id = t.employee_id
INNER JOIN project p on p.id = t.project_id
WHERE t.time_from >= NOW() - INTERVAL '1 month'
GROUP BY t.employee_id, e.name, p.name
ORDER BY e.name, p.name;
```

EXPLAIN ANALYZE result:

```
"QUERY PLAN"
"GroupAggregate (cost=86.74..99.25 rows=385 width=592) (actual time=0.170..0.178
rows=2 loops=1)"
"  Group Key: e.name, p.name, t.employee_id"
"  -> Sort (cost=86.74..87.70 rows=385 width=576) (actual time=0.123..0.125 rows=3
loops=1)"
"    Sort Key: e.name, p.name, t.employee_id"
"    Sort Method: quicksort Memory: 25kB"
"    -> Hash Join (cost=34.18..70.21 rows=385 width=576) (actual time=0.104..0.109
rows=3 loops=1)"
"      Hash Cond: (t.employee_id = e.id)"
"      -> Hash Join (cost=13.82..48.84 rows=385 width=438) (actual
time=0.072..0.076 rows=3 loops=1)"
"        Hash Cond: (t.project_id = p.id)"
"        -> Seq Scan on time_record t (cost=0.00..33.80 rows=453 width=24)
(actual time=0.024..0.026 rows=3 loops=1)"
"          Filter: (time_from >= (now() - '1 mon'::interval))"
```

```
"      -> Hash (cost=11.70..11.70 rows=170 width=426) (actual
time=0.035..0.035 rows=2 loops=1)"
"          Buckets: 1024 Batches: 1 Memory Usage: 9kB"
"          -> Seq Scan on project p (cost=0.00..11.70 rows=170 width=426)
(actual time=0.005..0.006 rows=2 loops=1)"
"          -> Hash (cost=14.60..14.60 rows=460 width=146) (actual time=0.005..0.006
rows=2 loops=1)"
"          Buckets: 1024 Batches: 1 Memory Usage: 9kB"
"          -> Seq Scan on employee e (cost=0.00..14.60 rows=460 width=146)
(actual time=0.004..0.004 rows=2 loops=1)"
"Planning Time: 0.619 ms"
"Execution Time: 0.311 ms"
```

After optimization:

- Uses INNER JOIN, so it uses hash join and scan only once on each project and employee table
- The query no longer use SubPlan
- The cost reduce from 7450.14..10514.78 into 86.74..99.25

Indexing strategies:

- Create index on time_from column in the time_record field, to speed up the interval filter
- Create index on employee_id and project_id in the time_record table for join